

CBF Programming Assignment

Overview

In this assignment, you will implement a content-based recommender as a LensKit recommender algorithm. LensKit provides tools to produce recommendations from a user; your task is to implement the logic of the recommender itself.

There are 2 parts to this assignment, implementing two variants of a TF-IDF recommender.

Downloads and Resources

- Project template (on course website)
- LensKit for Teaching website (links to relevant documentation and the LensKit tutorial video)
- JavaDoc for included code

Additionally, you will need:

- Java — download the Java 8 JDK. On Linux, install the OpenJDK ‘devel’ package (you will need the devel package to have the compiler).
- An IDE; I recommend IntelliJ IDEA Community Edition.

Notation

Here’s the mathematical notation we are using:

\vec{u} The user’s vector (in this assignment, the user profile vector).

\vec{i} The item vector.

$I(u)$ The set of items rated by user u .

u_t, i_t User u ’s or item i ’s score for tag t

r_{ui} User u ’s rating for item i .

μ_u The average of user u ’s ratings.

Part 1: TF-IDF Recommender with Unweighted Profiles (85 points)

Start by downloading the project template. This is a Gradle project; you can import it into your IDE directly (IntelliJ users can open the build.gradle file as a project). The code should compile as-is; you can test this by running the `build` Gradle target from your IDE, or running `./gradlew build` at the command line.

There are 3 things you need to implement to complete the first part of the assignment:

Compute item-tag vectors (the model) For this task, you need to modify the model builder (TFIDFModelBuilder, your modifications go in the get() method) to compute the unit-normalized TF-IDF vector for each movie in the data set. We provide the skeleton of this; TODO comments indicate where you need to implement missing pieces. When this piece is done, the model should contain a mapping of item IDs to TF-IDF vectors, normalized to unit vectors, for each item.

Build user profile for each query user The UserProfileBuilder interface defines classes that take a user's history – a list of ratings — and produce a vector representing that user's profile. For part 1, the profile should be the sum of the item-tag vectors of all items the user has rated positively (≥ 3.5 stars); this implementation goes in ThresholdUserProfileBuilder.

Generate item scores for each user The heart of the recommendation process in many LensKit recommenders is the score method of the item scorer, in this case TFIDFItemScorer. Modify this method to score each item by using cosine similarity: the score for an item is the cosine between that item's tag vector and the user's profile vector. Cosine similarity is defined as follows:

$$\cos(u, i) = \frac{\vec{u} \cdot \vec{i}}{\|\vec{u}\|_2 \|\vec{i}\|_2} = \frac{\sum_t u_t i_t}{\sqrt{\sum_t u_t^2} \sqrt{\sum_t i_t^2}}$$

You can run your program from the command line using Gradle:

```
...  
./gradlew recommendBasic -PuserId=42  
...
```

Try different user IDs.

Example Output for Unweighted User Profile

The following example gives actual outputs for users 42 and 91 in the data set. It was executed using ./gradlew recommendBasic -PuserId=42,91 in a Unix-like console.

```
recommendations for user 42:  
2081 (Little Mermaid, The (1989)): 0.500  
2671 (Notting Hill (1999)): 0.292  
2724 (Runaway Bride (1999)): 0.286  
1265 (Groundhog Day (1993)): 0.241  
58299 (Horton Hears a Who! (2008)): 0.224  
96861 (Taken 2 (2012)): 0.215  
1210 (Star Wars: Episode VI - Return of the Jedi (1983)): 0.187  
1202 (Withnail & I (1987)): 0.127  
26048 (Human Condition II, The (Ningen no joken II) (1959)): 0.110  
260 (Star Wars: Episode IV - A New Hope (1977)): 0.109
```

recommendations for user 91:

4816 (Zoolander (2001)): 0.343
8376 (Napoleon Dynamite (2004)): 0.252
1247 (Graduate, The (1967)): 0.232
4361 (Tootsie (1982)): 0.210
6350 (Laputa: Castle in the Sky (Tenkû no shiro Rapyuta) (1986)): 0.169
2167 (Blade (1998)): 0.161
2858 (American Beauty (1999)): 0.159
6947 (Master and Commander: The Far Side of the World (2003)): 0.158
5459 (Men in Black II (a.k.a. MIIB) (a.k.a. MIB 2) (2002)): 0.158
2724 (Runaway Bride (1999)): 0.151

Part 2: Weighted User Profile (15 points)

For this part, adapt your solution from Part 1 to compute weighted user profiles. Put your weighted user profile code in `WeightedUserProfileBuilder`.

In this variant, rather than just summing the vectors for all positively-rated items, compute a weighted sum of the item vectors for all rated items, with weights being based on the user's rating. Your solution should implement the following formula:

$$\vec{u} = \sum_{i \in I(u)} (r_{ui} - \mu_u) \vec{i}$$

Example Output for Weighted User Profile

The following example gives actual outputs for users 42 and 91 in the data set. It was executed using `./gradlew recommendWeighted -PuserId=42, 91` in a Unix-like console.

recommendations for user 42:

2081 (Little Mermaid, The (1989)): 0.191
2724 (Runaway Bride (1999)): 0.109
1265 (Groundhog Day (1993)): 0.107
2671 (Notting Hill (1999)): 0.093
96861 (Taken 2 (2012)): 0.082
58299 (Horton Hears a Who! (2008)): 0.077
1202 (Withnail & I (1987)): 0.056
1295 (Unbearable Lightness of Being, The (1988)): 0.047
7361 (Eternal Sunshine of the Spotless Mind (2004)): 0.039
3147 (Green Mile, The (1999)): 0.036

recommendations for user 91:

4816 (Zoolander (2001)): 0.438
8376 (Napoleon Dynamite (2004)): 0.323
1247 (Graduate, The (1967)): 0.296

4361 (Tootsie (1982)): 0.268
6350 (Laputa: Castle in the Sky (Tenkû no shiro Rapyuta) (1986)): 0.216
2167 (Blade (1998)): 0.206
6947 (Master and Commander: The Far Side of the World (2003)): 0.202
5459 (Men in Black II (a.k.a. MIIB) (a.k.a. MIB 2) (2002)): 0.202
2724 (Runaway Bride (1999)): 0.193
2161 (NeverEnding Story, The (1984)): 0.190

Submitting

Submit your code as a zip file to the TA (taijala@cs.umn.edu).

To create this zip file, please use the pre-created archive functionality in the Gradle build:

```
./gradlew prepareSubmission
```

This will ensure that your submission contains all required files. It will produce a submission file in `build/distributions`.